

SCI 4 OpenSHMEM TUTORIAL

Presenters: Tony Curtis, Deepak Eachempati, Dounia Khaldi, Aaron Welch

University of Houston, Texas


Oscar Hernandez, Pavel Shamis, Manju Gorentla Venkata

Oak Ridge National Laboratory

OpenSHMEM WORKSHOP OUTLINE

<http://www.openshmem.org/>

- Prerequisites
- Background
- Concepts
- History and Implementations
- The OpenSHMEM Project
- OpenSHMEM routines
- OpenSHMEM and Hardware
- Developing OpenSHMEM Applications
- OpenSHMEM Implementations
- OpenSHMEM: The Future...



OpenSHMEM

PREREQUISITES

- Knowledge of C/Fortran
- Familiarity with parallel computing
- Linux/UNIX command-line
- Useful for hands-on
 - 64-bit Linux (native, VM or remote)
 - E.g. Fedora, CentOS, Ubuntu, ...
 - Installation of GASNet, “fast” segment configuration preferable
 - <http://gasnet.lbl.gov/>
 - OpenSHMEM download, test-suite & demo programs
 - <https://github.com/openshmem-org/openshmem>

OpenSHMEM BACKGROUND (I)

- ▶ Large applications require lots of compute power
- ▶ Various approaches to providing this
 - ▶ Mainframe
 - ▶ SMP
 - ▶ Cluster
- ▶ All involve
 - ▶ Multiple things happening at once
 - ▶ ...Which needs...
 - ▶ Programming methods to
 - ▶ Express this
 - ▶ Take advantage of systems

OpenSHMEM

BACKGROUND (2)

- ▶ 2 main software paradigms
 - ▶ Threaded
 - ▶ Message-passing
- ▶ 2 main hardware paradigms
 - ▶ Single-image multiprocessor (SMP)
 - ▶ Distributed
 - ▶ Multiple machines with separate OS
 - ▶ Connected together

OpenSHMEM

BACKGROUND (3)

- ▶ Programming environments provide abstraction
- ▶ In particular
 - ▶ A language or library can be used on many machine types
 - ▶ Implementation hides differences & leverages features
- ▶ 2 dominant models
 - ▶ MPI
 - ▶ OpenMP
- ▶ First, a little background for context...

OpenSHMEM BACKGROUND (4)

- ▶ Concurrent
 - ▶ Multiple things logically happen at once
 - ▶ May be emulated
 - ▶ E.g. time slicing on shared machine
- ▶ Parallel
 - ▶ = Concurrent +
 - ▶ Things really happen independently
 - ▶ On separate processors
- ▶ Work is partitioned in some way across resources

OpenSHMEM

BACKGROUND (5)

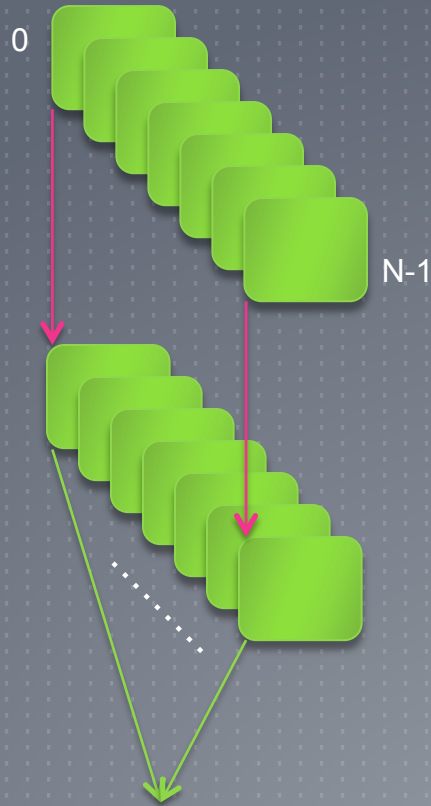
- ▶ Different ways of partitioning work
 - ▶ different tasks * different data
- ▶ MPMD = multiple program, multiple data
- ▶ SPMD = single program, multiple data
- ▶ SPSD = single program, single data
 - ▶ Just good old sequential

<http://en.wikipedia.org/wiki/SPMD>

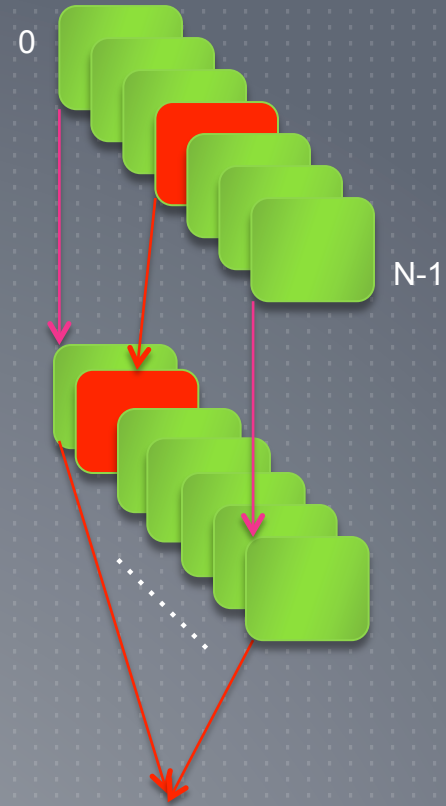
OpenSHMEM BACKGROUND (6)

- ▶ SPMD
 - ▶ Program launches many processes
 - ▶ Each starts with same code (SP)
 - ▶ And then typically operates on some specific part of the data (MD)
 - ▶ Processes may then communicate with each other
 - ▶ Share common data
 - ▶ Broadcast work
 - ▶ Collect results
- ▶ PVM and MPI are well-known examples
- ▶ OpenSHMEM is SPMD

OpenSHMEM BACKGROUND (7)



Independent execution



Communication/Synchronization

OpenSHMEM

BACKGROUND (8)

- ▶ Address Spaces

- ▶ Global vs. distributed
- ▶ OpenMP has global (shared) space
- ▶ MPI has partitioned space
 - ▶ Private data exchanged via messages
- ▶ OpenSHMEM is “partitioned global address space”
 - ▶ PGAS
- ▶ Has private *and* shared data
 - ▶ Shared data accessible directly by other processes

OpenSHMEM

BACKGROUND (9)

- ▶ The PGAS family
 - ▶ Libraries include...
 - ▶ GASNet
 - ▶ ARMCI / Global Arrays
 - ▶ CCI
 - ▶ GASPI/GPI
 - ▶ Languages include...
 - ▶ Chapel
 - ▶ Titanium
 - ▶ X10
 - ▶ UPC
 - ▶ CAF
 - ▶ (Often built on these libraries)

OpenSHMEM

BACKGROUND (10)

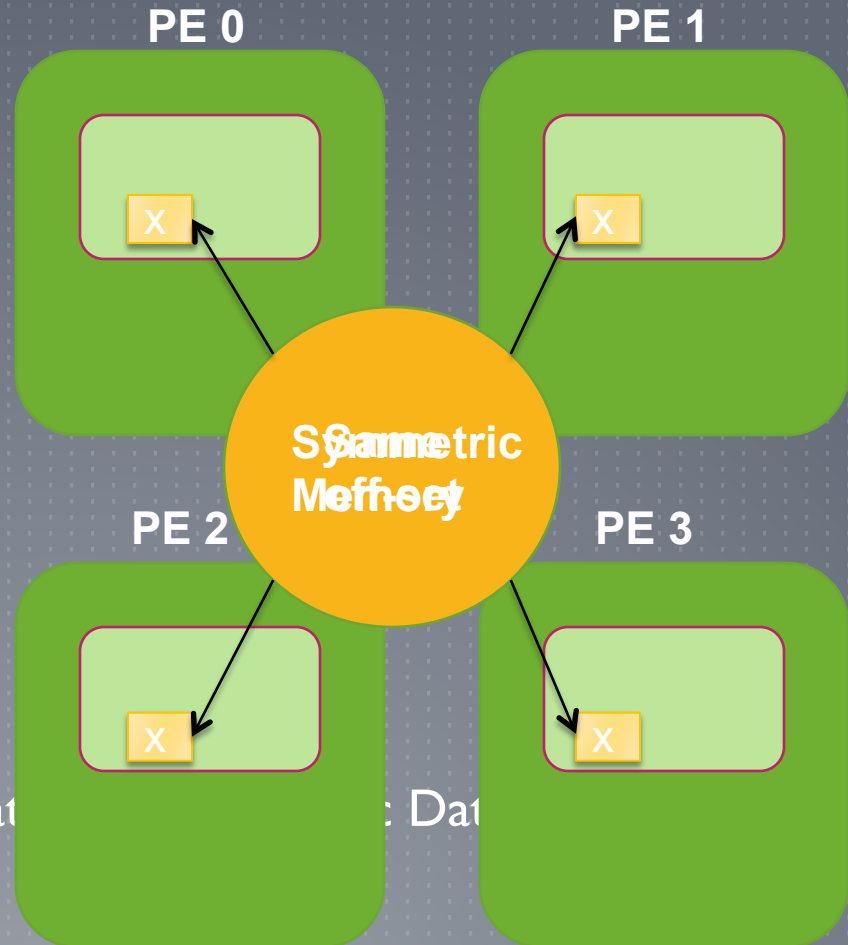
- ▶ OpenSHMEM is SPMD parallel programming library
 - ▶ Library of functions similar in feel to using MPI (e.g. *shmem_get()*)
- ▶ Available for C, C++ and Fortran
- ▶ Used for programs that
 - perform computations in separate address spaces and
 - explicitly pass data to and from different processes in the program.
- ▶ The processes participating in shared memory applications are referred to as processing elements (PEs).
- ▶ OpenSHMEM routines supply remote data transfer, work-shared broadcast and reduction, barrier synchronization, and atomic memory operations.

OpenSHMEM CONCEPTS (I)

- ▶ Symmetric Variables
 - ▶ Arrays or variables that exist with the **same size, type, and relative address** on all PEs.
 - ▶ The following kinds of data objects are symmetric:
 - ▶ Fortran data objects in common blocks or with the **SAVE** attribute.
 - ▶ Non-stack C and C++ variables.
 - ▶ Fortran arrays allocated with **shpalloc**
 - ▶ C and C++ data allocated by **shmalloc**

OPENSHPMEM CONCEPTS (2)

```
#include <shmem.h>
int main (void)
{
    int *x;
    ...
    start_pes(0);
    ...
    x = (int*) shmalloc(sizeof(x));
    ...
    shmem_barrier_all();
    ...
    shfree(x);
    return 0;
}
```



OpenSHMEM

HISTORY AND IMPLEMENTATIONS

- ▶ Cray
 - ▶ SHMEM first introduced by Cray Research Inc. in 1993 for Cray T3D
 - ▶ Platforms: Cray T3D, T3E, PVP, XT series
- ▶ SGI
 - ▶ Owns the “rights” for SHMEM
 - ▶ Baseline for OpenSHMEM development (Altix)
- ▶ Quadrics (company out of business)
 - ▶ Optimized API for QsNet
 - ▶ Platform: Linux cluster with QsNet interconnect
- ▶ Others
 - ▶ HP SHMEM, IBM SHMEM
 - ▶ GPShMEM (cluster with ARMCI & MPI support, old)

Note: SHMEM was not defined by any one standard.

OpenSHMEM

DIVERGENT IMPLEMENTATIONS (I)

- ▶ Many forms of initialization
 - ▶ Include header shmem.h to access the library
 - ▶ E.g. `#include <shmem.h>` , `#include <mpp/shmem.h>`
 - ▶ `start_pes`, `shmem_init`: Initializes the calling PE
 - ▶ `my_pe`: Get the PE ID of local processor
 - ▶ `num_pes`: Get the total number of PEs in the system

SGI		Quadrics	Cray	
Fortran	C/C++	C/C++	Fortran	C/C++
<code>start_pes</code>	<code>start_pes(0)</code>	<code>shmem_init</code>	<code>start_pes</code>	<code>start_pes</code>
			<code>shmem_init</code>	<code>shmem_init</code>
<code>shmem_my_pe</code>	<code>shmem_my_pe</code>		<code>shmem_my_pe</code>	<code>shmem_my_pe</code>
<code>shmem_n_pes</code>	<code>shmem_n_pes</code>		<code>shmem_n_pes</code>	<code>shmem_n_pes</code>
<code>NUM_PES</code>	<code>num_pes</code>	<code>num_pes</code>	<code>NUM_PES</code>	
<code>MY_PE</code>	<code>my_pe</code>	<code>my_pe</code>		

OpenSHMEM

DIVERGENT IMPLEMENTATIONS (2)

Hello World (SGI on Altix)

```
#include <stdio.h>
#include <mpp/shmem.h>

int main(void)
{
    int me, npes;
    start_pes(0);
    npes = _num_pes();
    me = _my_pe();
    printf("Hello from %d of %d\n", me, npes);
    return 0;
}
```

Hello World (SiCortex)

```
#include <stdio.h>
#include <shmem.h>

int main(void)
{
    int me, npes;
    shmem_init();
    npes = num_pes();
    me = my_pe();
    printf("Hello from %d of %d\n", me, npes);
    return 0;
}
```

OpenSHMEM

THE PROJECT

- ▶ <http://www.openshmem.org/>
- ▶ Standardized specification
- ▶ Reference Library of spec.
- ▶ Tutorials & other educational material
- ▶ Vendor products & information
- ▶ Community involvement, talk to each other!
- ▶ Tool-chain ecosystem

OpenSHMEM ROUTINES

- ▶ **Initialization and Program Query**
- ▶ **Data transfers**
- ▶ **Synchronization mechanisms**
- ▶ **Collective communication**
- ▶ **Atomic Memory Operations**
- ▶ **Address Manipulation, Data Cache control**
 - ▶ Not supported by all SHMEM implementations

OpenSHMEM INITIALIZATION & QUERY

- ▶ **void start_pes(int n)**
 - ▶ Initialize the OpenSHMEM program
 - ▶ “n” means “number of PEs” but now ignored, set to 0
 - ▶ Number of PEs taken from invoking environment
 - ▶ E.g. from MPI or job scheduler
 - ▶ PEs numbered 0 .. (N – 1) in flat space
- ▶ **int _num_pes(void)**
- ▶ **int shmem_n_pes(void)**
 - ▶ **return number of PEs in this program**
- ▶ **int _my_pe(void)**
- ▶ **int shmem_my_pe(void)**
 - ▶ **return “rank” of calling PE**

OpenSHMEM

DATA TRANSFER (I)

► Put

- Single variable
 - **void shmem_TYPE_p(TYPE *target, TYPE value, int pe)**
 - TYPE = double, float, int, long, short
- Contiguous object
 - **void shmem_TYPE_put(TYPE *target, const TYPE *source, size_t nelems, int pe)**
 - For C: TYPE = double, float, int, long, longdouble, longlong, short
 - For Fortran: TYPE = complex, integer, real, character, logical
 - **void shmem_putSS(void *target, const void *source, size_t nelems, int pe)**
 - Storage Size (SS, bits) = 32, 64, 128, mem (any size)
- Target must be symmetric

OpenSHMEM

DATA TRANSFER (2)

- ▶ Example: Cyclic communication via puts

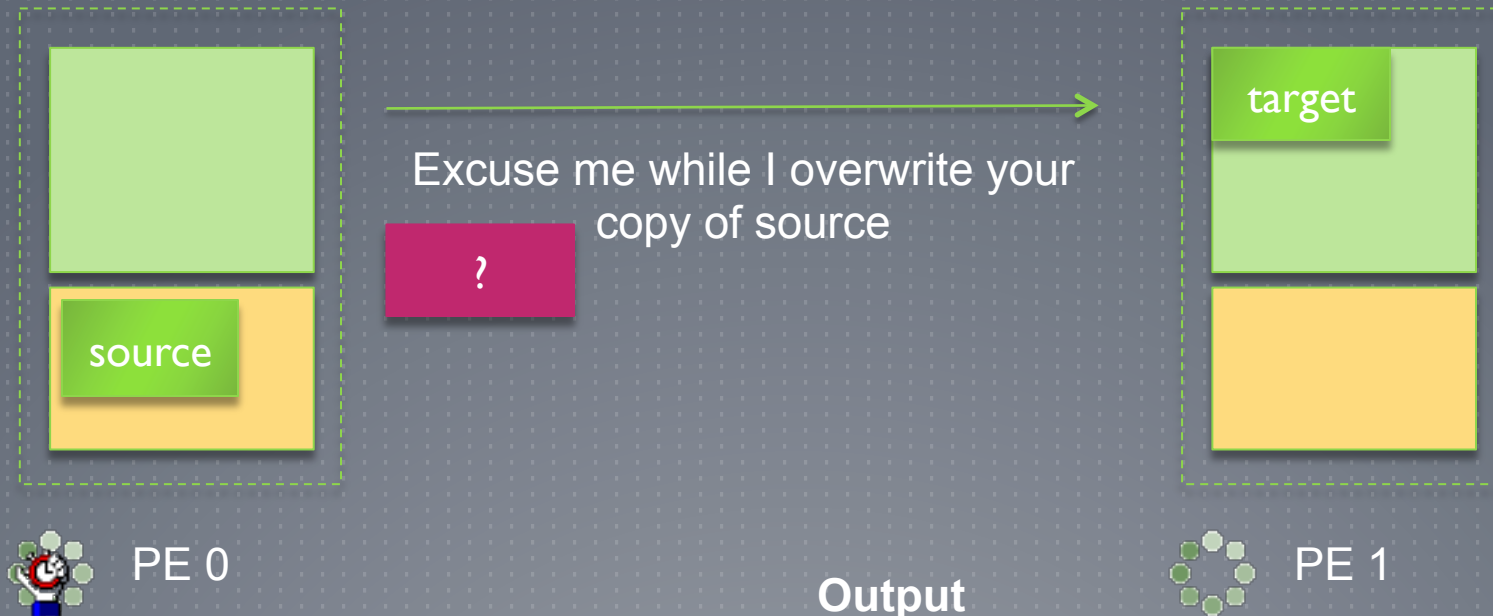
```
{  
    /*Initializations*/  
    int src;  
    int *dest;  
    ....  
    start_pes(0);  
    ....  
    src = me; Automatic data element  
    dest = (int *) shmalloc (sizeof (*dest)); Symmetric data element  
    nextpe = (me + 1) % npes; /*wrap around*/  
  
    shmem_int_put (dest, &src, 1, nextpe);  
    more_work_goes_here (... Synchronization before use  
    shmem_barrier_all();  
    x = dest * 0.995 + 45 * y;  
    ...  
}
```

Points To Remember

- 'Destination' has to be symmetric
- Consecutive puts are not guaranteed to finish in order
- Put returns after the data has been copied out of the source
- Completion guaranteed only after synchronization

OpenSHMEM

DATA TRANSFER (3): PUT



Output

target[0] on PE 1 is 1
target[1] on PE 1 is 2
target[2] on PE 1 is 3
target[3] on PE 1 is 4
...
target[9] on PE 1 is 10

Shared Address Space

Private Address Space

OpenSHMEM

DATA TRANSFER (4)

► Get

- Single variable
 - **TYPE shmem_TYPE_g(TYPE *target, TYPE value, int pe)**
 - For C: TYPE = double, float, int, long, longdouble, longlong, short
 - For Fortran: TYPE=complex, integer, real, character, logical
- Contiguous object
 - **void shmem_TYPE_get(TYPE *target, const TYPE *source, size_t nelems, int pe)**
 - For C: TYPE = double, float, int, long, longdouble, longlong, short
 - For Fortran: TYPE=complex, integer, real, character, logical
 - **void shmem_getSS(void *target, const void *source, size_t nelems, int pe)**
 - Storage Size (SS, bits) = 32, 64, 128, mem (any size)
- Source must be symmetric

OpenSHMEM DATA TRANSFER (5)

- ▶ Example: Summation at PE 0

```
{  
    /*Initializations*/  
    int *src;  
    int target, sum;  
    ....  
    start_pes(0);  
    ....  
    src = (int *) shmalloc (sizeof (*src));  
    src = me;  
    sum=me;  
    if(me == 0){  
        for(int i = 1,i<num_pes();i++){  
            shmem_int_get(&target, src, i  
            sum = sum + target;  
        }  
    }  
    ...  
}
```

Automatic data element

Symmetric data element

No synchronization before use

Points To Remember

- 'Source' has to be remotely accessible
- Consecutive gets finish in order
- The routines return after the data has been delivered to the 'target' on the local PE

OpenSHMEM

DATA TRANSFER (6)

► Strided put/get

► **void shmem_TYPE_iput(TYPE *target, const TYPE *source, ptrdiff_t tst, ptrdiff_t sst, size_t nelems, int pe)**

- For C: TYPE = double, float, int, long, longdouble, longlong, short
- For Fortran: TYPE = complex, integer, real, character, logical
- tst and sst indicate stride between accesses of target and source resp.

► And the sized variants as for put/get

OpenSHMEM

DATA TRANSFER (7)

► Put vs. Get

- Put call completes when data is “being sent”
- Get call completes when data is “stored locally”
- Cannot assume put has written until later synchronization
 - Data still in transit
 - Partially written at target
 - Put order changed by e.g. network
- Puts allow overlap
 - Communicate
 - Compute
 - Synchronize

OpenSHMEM SYNCHRONIZATION (I)

- ▶ Active Sets
 - ▶ Way to specify a subset of PEs
 - ▶ A triple:
 - ▶ Start PE
 - ▶ Stride (\log_2)
 - ▶ Size of set
- ▶ Limitations
 - ▶ Stride must be powers of 2
 - ▶ Only define 'regular' PE sub-groups

OpenSHMEM SYNCHRONIZATION (2)

► Quick look at Active Sets

► Example 1

► $PE_start = 0, \log PE_stride = 0, PE_size = 4$

ACTIVE SET? PE 0, PE 1, PE 2, PE 3

► Example 2

► $PE_start = 0, \log PE_stride = 1, PE_size = 4$

ACTIVE SET? PE 0, PE 2, PE 4, PE 6

► Example 3

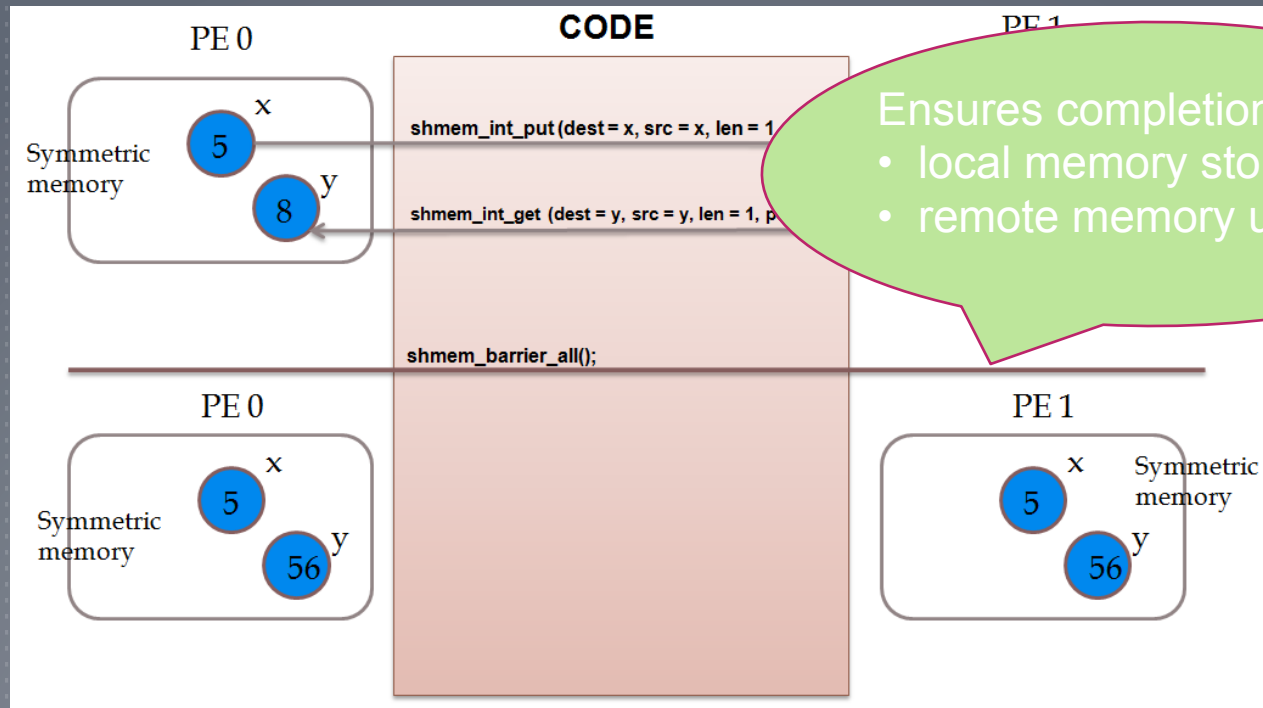
► $PE_start = 2, \log PE_stride = 2, PE_size = 3$

ACTIVE SET? PE 2, PE 6, PE 10

OpenSHMEM SYNCHRONIZATION (3)

- ▶ Barrier (Group synchronization)
 - ▶ `void shmem_barrier_all()`
 - ▶ Suspend PE execution until all PEs call this function
 - ▶ `void shmem_barrier(int PE_start, int PE_stride, int PE_size, long *pSync)`
 - ▶ Barrier operation on subset of PEs
- ▶ *pSync* is a symmetric work array that allows different barriers to operate simultaneously

OpenSHMEM SYNCHRONIZATION (4)



`shmem_barrier_all()` synchronizes all executing PEs

OpenSHMEM

SYNCHRONIZATION (5)

- ▶ Conditional wait (P2P synchronization)
 - ▶ Suspend until local symmetric variable NOT equal to the value specified
 - ▶ **void shmem_wait(long *var, long value)**
 - ▶ **void shmem_TYPE_wait(TYPE *var, TYPE value)**
 - ▶ For C: TYPE = int, long, longdouble, longlong, short
 - ▶ For Fortran: TYPE = complex, integer, real, character, logical
- ▶ Specific conditional wait
 - ▶ Similar to the generic wait except the comparison can now be
 - ▶ \geq , $>$, $=$, \neq , $<$, \leq
 - ▶ **void shmem_wait_until(long *var, int cond, long value)**
 - ▶ **void shmem_TYPE_wait_until(TYPE *var, int cond, TYPE value)**
 - ▶ TYPE = int, long, longlong, short

OpenSHMEM

SYNCHRONIZATION (6)

Fence

Ordering of outgoing write (put) operations to a single PE

`void shmem_fence()`

Quiet

Ordering of all outgoing puts from the calling PE (on some implementations;
fence = quiet)

`void shmem_quiet()`

OpenSHMEM SYNCHRONIZATION (7)

Example Fence

```
int main (int argc, char **argv)
{
    ....
    ....
    shmem_int_put (dest1, src1, 1, nexttpe);
    shmem_fence();
    shmem_int_put (dest1, src2, 1, nexttpe);
    ....
    shmem_barrier_all ();
    shfree (dest);

    return 0;
}
```

Example Quiet

```
int main (int argc, char **argv)
{
    ....
    ....
    shmem_int_put (dest1, src1, 1, nexttpe);
    shmem_int_put (dest2, src2, 1, nexttpe+1);
    shmem_quiet();
    ....
    shmem_barrier_all ();
    shfree (dest);

    return 0;
}
```

OpenSHMEM

COLLECTIVE COMMUNICATION (I)

► Broadcast

- One-to-all symmetric communication
- No update on root

► `void shmem_broadcastSS(void *target, void *source, size_t nelems, int PE_root, int PE_start, int PE_stride, int PE_size, long *pSync)`

Storage Size (SS, bits) = 32, 64

OpenSHMEM

COLLECTIVE COMMUNICATION (2)

37

```
...  
...  
int *target, *source;  
target= (int *) shmalloc( sizeof(int) );  
source= (int *) shmalloc( sizeof(int) );  
*target= 0;  
*source= 101;  
if (me == 1) {  
    *source = 222;  
}  
shmem_barrier_all();  
shmem_broadcast32(target, source, 1 0, 0, 0, 4, pSync);  
  
printf("target on PE %d is %d\n", _my_pe(), *target);  
...
```

Output

target on PE 0 is 0
target on PE 1 is 222
target on PE 2 is 222
target on PE 3 is 222

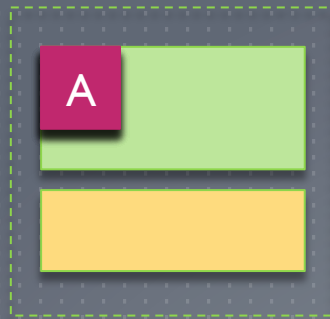
collective operation

of the PE
active set

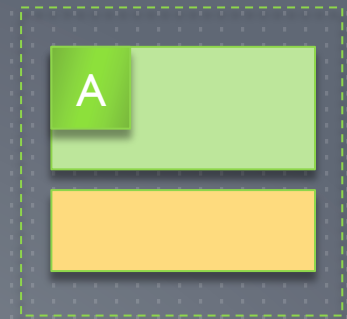
Code snippet showing working of shmem_broadcast

OpenSHMEM


COLLECTIVE COMMUNICATION (3)




PE 0



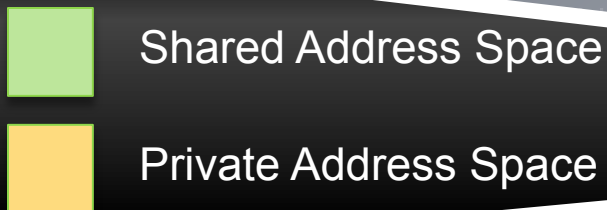
PE 3



PE 1



PE 2



OpenSHMEM

COLLECTIVE COMMUNICATION (4)


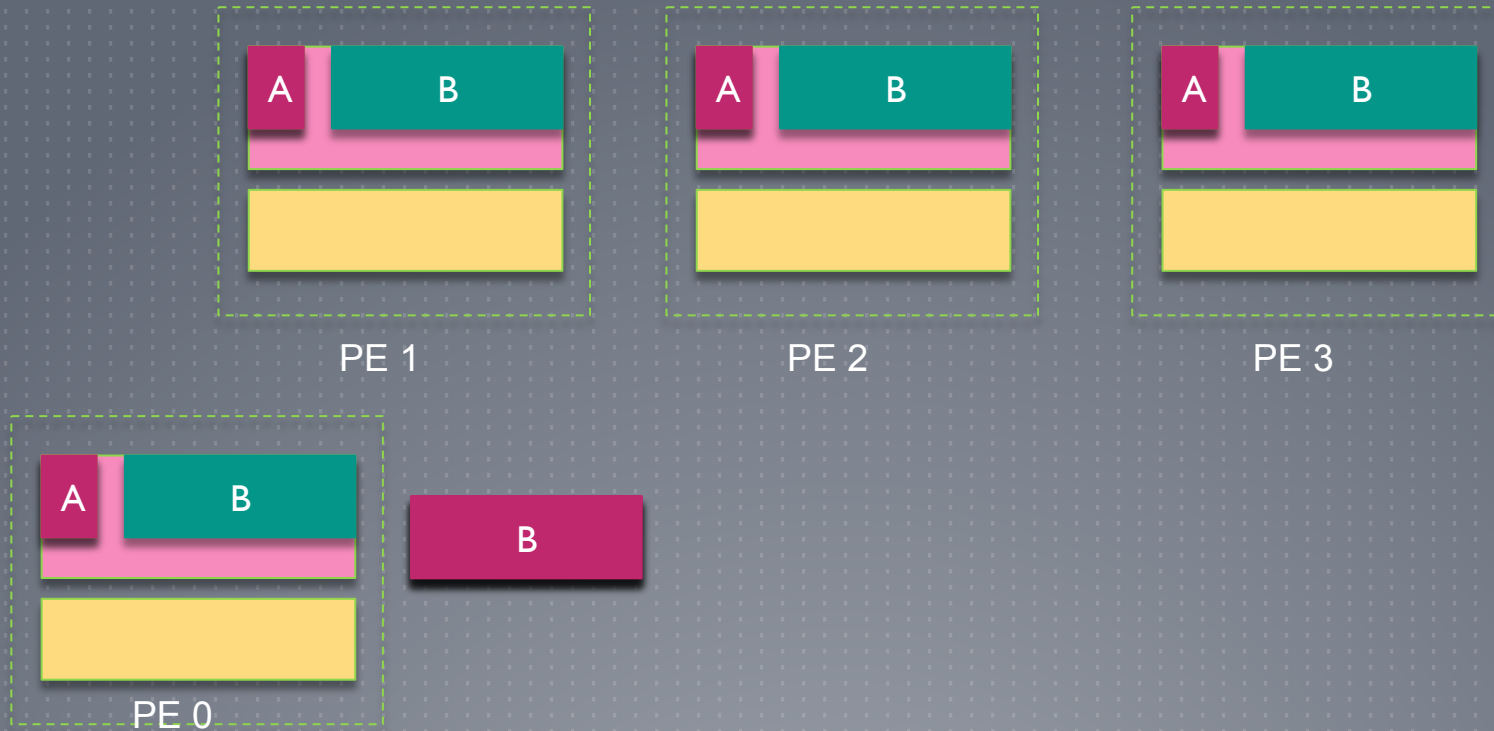
Storage Size (SS, bits) = 32, 64

► Collection

- Concatenates blocks of symmetric data from multiple PEs to an array in every PE
- Each PE can contribute different amounts
- **void shmem_collectSS(void *target, void *source, size_t nelems, int PE_start, int PE_stride, int PE_size, long *pSync)**
- Concatenation written on all participating PEs
- **shmem_fcollect variant**
 - When all PEs contribute exactly same amount of data
 - PEs know exactly where to write data, so no offset lookup overhead

OpenSHMEM

COLLECTIVE COMMUNICATION (5)



OpenSHMEM

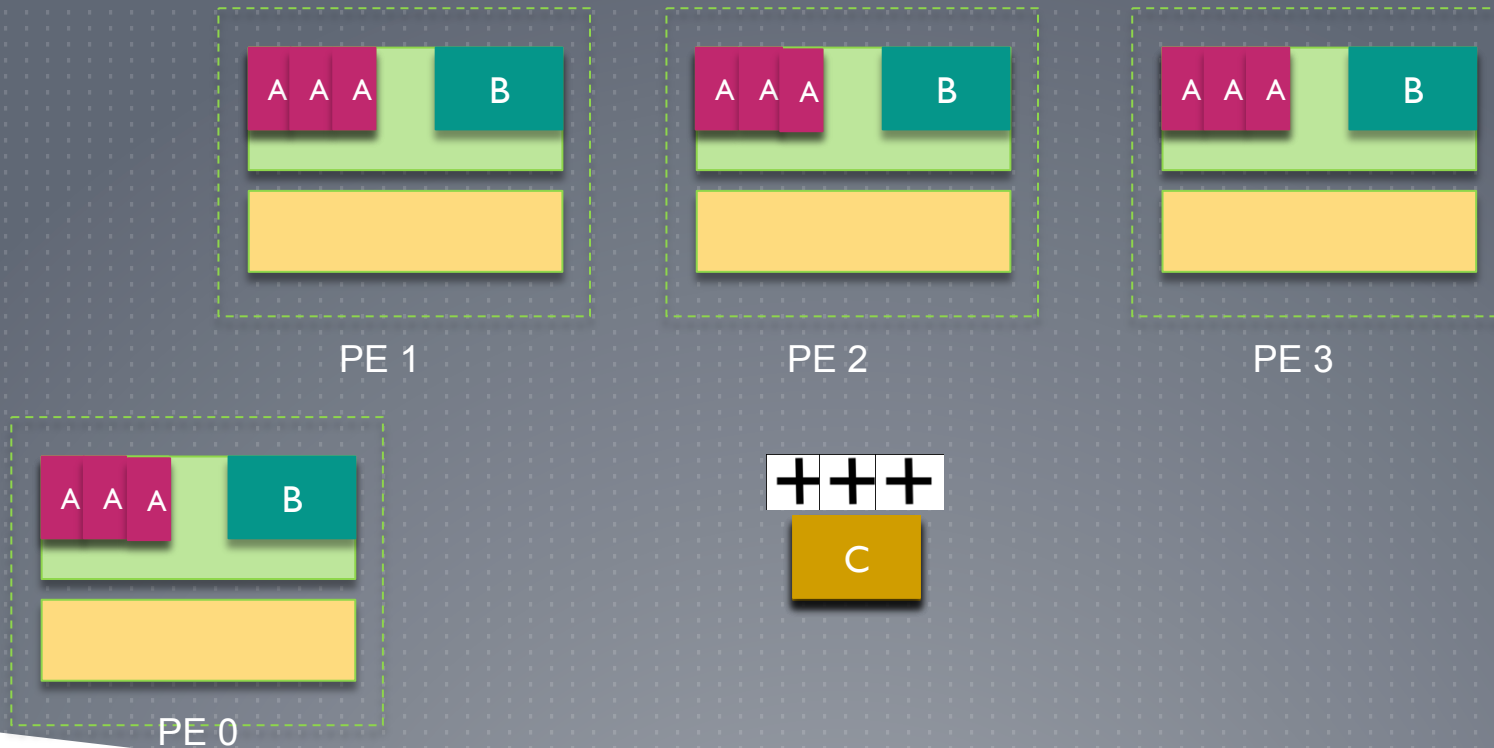
COLLECTIVE COMMUNICATION (6)

► Reductions

- Perform commutative operation across symmetric data set
 - `void shmem_TYPE_OP_to_all(TYPE *target, TYPE *source, int nreduce, int PE_start, int PE_stride, int PE_size, TYPE *pWrk, long *pSync)`
 - Logical OP = and, or, xor
 - Extrema OP = max, min
 - Arithmetic OP = prod(uct), sum
 - TYPE = int, long, longlong, longdouble, short, complex
- Reduction performed and stored on all participating PEs
- pWrk and pSync allow interleaving
- E.g. compute arithmetic mean across set of PEs
 - `sum_to_all / PE_size`

OpenSHMEM

COLLECTIVE COMMUNICATION (7)



Shared Address Space

Private Address Space

OpenSHMEM

ATOMIC OPERATIONS (I)

- ▶ What does “atomic” mean anyway?
 - ▶ Indivisible operation on symmetric variable
 - ▶ No other operation can interpose during update
- ▶ But “no other operation” actually means...?
 - ▶ No other atomic operation
 - ▶ Can't do anything about other mechanisms interfering
 - ▶ E.g. thread outside of OpenSHMEM program
 - ▶ Non-atomic OpenSHMEM operation
- ▶ Why this restriction?
 - ▶ Implementation in hardware

OpenSHMEM

ATOMIC OPERATIONS (2)

▶ Atomic Swap

▶ Unconditional

- ▶ **long shmem_swap(long *target, long value, int pe)**
- ▶ **TYPE shmem_TYPE_swap(TYPE *target, TYPE value, int pe)**
 - ▶ TYPE = double, float, int, long, longlong
 - ▶ Return old value from symmetric target

▶ Conditional

- ▶ **TYPE shmem_TYPE_cswap(TYPE *target, TYPE cond, TYPE value, int pe)**
 - ▶ TYPE = int, long, longlong
- ▶ Only if “cond” matches value on target

OpenSHMEM

ATOMIC OPERATIONS (3)

► Arithmetic

- increment (= add 1) & add value
- `void shmem_TYPE_inc(TYPE *target, int pe)`
- `void shmem_TYPE_add(TYPE *target, TYPE value, int pe)`
 - TYPE = int, long, longlong
- Fetch-and-increment & fetch-and-add value
- `TYPE shmem_TYPE_finc(TYPE *target, int pe)`
- `TYPE shmem_TYPE_fadd(TYPE *target, TYPE value, int pe)`
 - TYPE = int, long, longlong
- Return previous value at target on PE

OpenSHMEM

ATOMIC OPERATIONS (4)

```
...  
...  
long *dest;  
dest = (long *) shmalloc( sizeof(*dest) );  
*dest= me;  
shmem_barrier_all();  
....  
new_val = me;  
if (me== 1) {  
    swapped_val = shmem_long_swap(target, new_val, 0);  
    printf("%d: target = %d, swapped = %d\n", me, *target, swapped_val);  
}  
shmem_barrier_all();  
...
```

OpenSHMEM

ATOMIC OPERATIONS (5)

► Locks

- Symmetric variables
- Acquired and released to define mutual-exclusion execution regions
 - Only 1 PE can enter at a time
- **void shmem_set_lock(long *lock)**
- **void shmem_clear_lock(long *lock)**
- **int shmem_test_lock(long *lock)**
 - Acquire lock if possible, return whether or not acquired
 - But don't block...
- Initialize lock to 0. After that managed by above API
- Can be used for updating distributed data structures

OpenSHMEM ACCESSIBILITY

- ▶ `int shmem_pe_accessible(int pe)`
 - ▶ Can this PE talk to the given PE?
- ▶ `int shmem_addr_accessible(void *addr, int pe)`
 - ▶ Can this PE address the named memory location on the given PE?
- ▶ In SGI SHMEM used for mixed-mode MPI/SHMEM programs
 - ▶ In “pure” OpenSHMEM, could just return “1”
- ▶ Could in future be adapted for fault-tolerance

OpenSHMEM

ADDRESSES & CACHE (I)

▶ Address manipulation


- ▶ `void *shmem_ptr(void *addr, int pe)`
 - ▶ Returns a pointer to a data object on a remote PE
- ▶ Only of use on platforms where memory physically accessible
- ▶ i.e. puts/gets are simple memory accesses

OpenSHMEM

ADDRESSES & CACHE (2)

► Cache control

- `shmem_clear_cache` - Disables automatic cache coherency mode
- `shmem_set_cache` - Enables automatic cache coherency mode
- `shmem_set_cache_inv` - Enables automatic line cache coherency mode
- `shmem_udcflush` - Flushes the entire user data cache coherency
- `shmem_udcflush` - Makes coherent all user data



OpenSHMEM

HARDWARE (I)

- ▶ Where is OpenSHMEM used?
 - ▶ Mainly clusters these days
 - ▶ Infiniband and similar networks
 - ▶ Why?
- ▶ Remote direct memory access (RDMA)
 - ▶ Network hardware writes directly into registered region of process memory
 - ▶ Without interrupting remote process(or)
 - ▶ Put symmetric memory areas here

Infiniband
Myrinet
Quadrics
SeaStar
RoCE

OpenSHMEM

HARDWARE (2)

- ▶ Offload
 - ▶ Infiniband HCAs can do
 - ▶ Atomics
 - ▶ Collectives
 - ▶ Memory pinning
 - ▶ Meaning CPU free to do other things
 - ▶ Reduced software footprint (QPs)
 - ▶ OpenSHMEM library issues offload instructions rather than doing atomics etc.

Developing OpenSHMEM Applications

OpenSHMEM

LOOKING FOR OVERLAPS (I)

- ▶ How to identify overlap opportunities
 - ▶ Put is not an indivisible operation
 - ▶ Send local, reuse local, on-wire, stored
 - ▶ Can do useful work on other data in between

OpenSHMEM

LOOKING FOR OVERLAPS (2)

- ▶ How to identify overlap opportunities
 - ▶ General principle:
 - ▶ Identify independent tasks/data
 - ▶ Initiate action as early as possible
 - ▶ Put/barrier/collective
 - ▶ Interpose independent work
 - ▶ Synchronize as late as possible

OpenSHMEM

LOOKING FOR OVERLAPS (3)

- ▶ How to identify overlap opportunities
 - ▶ How could we change OpenSHMEM to get even more overlap?
 - ▶ Divide application into distinct communication and computation phases to minimize synchronization points
 - ▶ Use of point-to-point synchronization as opposed to collective synchronization

OPENSHMEM

LOOKING FOR OVERLAPS (4)

- ▶ How to identify overlap opportunities
 - ▶ Shmalloc
 - ▶ Size check, allocate, barrier_all
 - ▶ Opportunities to do other work after local allocation
 - ▶ Then wait in barrier later
 - ▶ Return handle for synch.

OPENSHMEM

LOOKING FOR OVERLAPS (5)

- ▶ How to identify overlap opportunities
 - ▶ “_nb” put/get calls
 - ▶ Local data not free for reuse on return
 - ▶ Return handle for later synch.

SOME OpenSHMEM IMPLEMENTATIONS

- ▶ Reference Library: University of Houston
 - ▶ On top of GASNet for portability
 - ▶ <http://www.openshmem.org/>
- ▶ ScalableSHMEM: Mellanox
 - ▶ For Mellanox Infiniband solutions
 - ▶ <http://www.mellanox.com/products/shmem>
- ▶ Portals-SHMEM: open-source
 - ▶ For Portals clusters
 - ▶ <http://code.google.com/p/portals-shmem/>
- ▶ Open-MPI
 - ▶ <http://www.open-mpi.org/>

OpenSHMEM

SUMMARY

- ▶ SPMD Library for C and Fortran programs
- ▶ Point-to-point data transfer
- ▶ Broadcast/collective transfer operations
- ▶ Synchronization
- ▶ Atomic operations

OpenSHMEM

REFERENCES

- ▶ Stephen W. Poole, Oscar Hernandez, Jeffery A. Kuehn, Galen M. Shipman, Anthony Curtis, and Karl Feind. OpenSHMEM - Toward a Unified RMA Model. Published in Encyclopedia of Parallel Computing, Springer US. Pages 1379-1391, 2011
- ▶ OpenSHMEM and Related Technologies. Experiences, Implementations, and Tools, First Workshop, OpenSHMEM 2014, Annapolis, MD, USA, March 4-6, 2014, Proceedings
- ▶ OpenSHMEM & Infiniband Research: visit DK Panda's group
 - ▶ <http://nowlab.cse.ohio-state.edu/>

ACKNOWLEDGEMENTS



This work was supported by the United States Department of Defense & used resources of the Extreme Scale Systems Center at Oak Ridge National Laboratory.